

# EVM

# Lesson Plan

Ethereum storage

Ethereum accounts

Ethereum transactions

EVM languages

Ethereum clients and mining

Connecting to a test network

# References

[DEVCON1: Understanding the Ethereum Blockchain Protocol - Vitalik Buterin](#)

[Mastering Ethereum](#) by Andreas Antonopoulos

[White paper](#)


[Beige Paper](#)

[Yellow Paper](#)

[EVM languages](#)

## Ethereum Data Structures

Ethereum uses Merkle Patricia Tries / Radix Tries for their searching performance and low memory footprint





# Data Structures

## There are 3 Tries


- World State
- Transaction
- Transaction Receipt

See : [Ethereum State Trie Architecture Explained](#)

# World and Account State

## Account storage contents Trie

A mapping between integer keys (KEC) and integer values (RLP)



## Account, $\sigma[\text{address}]$

RLP data structure

### nonce, $\sigma[\text{address}]_n$

scalar; the number of transactions sent from this address or, in the case of accounts with associated code, the number of contract-creations made by this account.

### balance, $\sigma[\text{address}]_b$

scalar; the number of Wei owned by this address

### storageRoot, $\sigma[\text{address}]_s$


256-bit hash of the root node of a Merkle Patricia tree that encodes the storage contents of the account (a mapping between 256-bit integer values), encoded into the trie as a mapping from the Keccak 256-bit hash of the 256-bit integer keys to the RLP-encoded 256-bit integer

### codeHash, $\sigma[\text{address}]_c$

Hash of the EVM code of this account - the code that gets executed should this address receive a message call; it is immutable and thus, unlike all other fields, cannot be changed after construction. All such code fragments are contained in the state database under their corresponding hashes for later retrieval.

## World State Trie, $\sigma$

A mapping between addresses and account states. Stored as a merkle-paritica tree



# Transactions

## Transaction, $T$

**nonce,  $T_n$**

Scalar; the number of transactions sent by the sender

**gasPrice,  $T_p$**

scalar; the number of Wei to be paid per unit of gas for all computation costs incurred as a result of execution of this transaction

**gasLimit,  $T_g$**

scalar; the max amount of gas that should be used in executing this transaction. Paid before any computation is done, may not be increased later.

**to,  $T_t$**


160-bit address of the message call's recipient or, for a contract creation transaction, 0 (zero bytes).

**value,  $T_v$**

scalar; the number of Wei to be transferred to the message call's recipient or, in the case of contract creation, as an endowment to the newly created account.

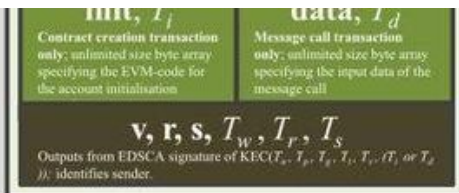
## Transaction Trie, $T$

A merkle-paritica tree of transactions to include

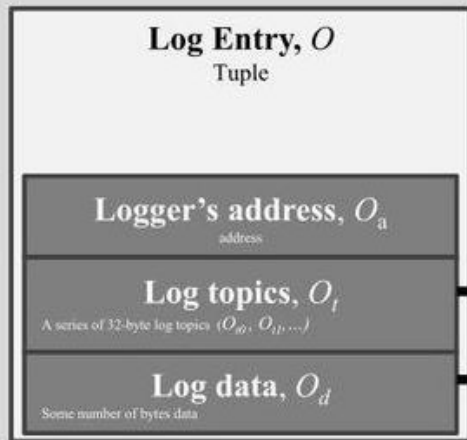
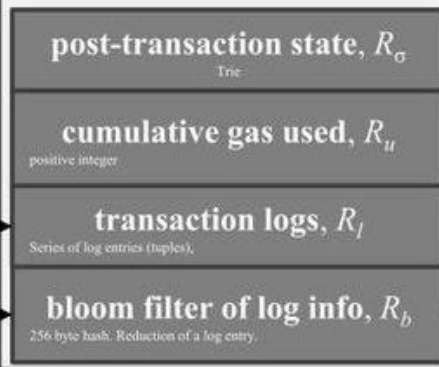




# Transaction Receipts and Logs




**Transaction Receipt,  $B_R[i]$**   
*( $i$ =transaction no)*  
 Tuple



## Transaction Receipts Trie

Index keyed trie



## Transaction and Transaction Receipt Tries

### Purpose:

- **Transaction Tries**: records transaction request vectors
- **Transaction Receipt Tries**: records the transaction outcome

### Content:

Parameters used in composing a **Transaction Trie** [details in section 4.3 of the [yellow paper](#)]:

- nonce,
- gas price,
- gas limit,
- recipient,
- transfer value,
- transaction signature values, and
- account initialization (if transaction is of contract creation type), or transaction data (if transaction is a message call)

Parameters used in composing a **Transaction Receipt Trie** [details in section 4.4.1 of the [yellow paper](#)]:

- post-transaction state,
- the cumulative gas used,
- the set of logs created through execution of the transaction, and
- the Bloom filter composed from information in those logs


# Ethereum Block Fields

- Header
  - Parent Hash
  - Ommers Hash
  - Beneficiary
  - State Root
  - Transaction Root
  - Receipts Root
  - Logs Bloom
  - Difficulty
  - Block Number
  - Gas Limit
  - Gas Used
  - Timestamp
  - Extra Data
  - Mix Hash
  - Nonce
- Ommer Block Headers
- Transactions

See [Beige Paper](#)

# Ommers (Uncle) Blocks

Unlike Bitcoin, Ethereum gives rewards to blocks that have been produced and are valid, but do not form part of the canonical chain. Has to be within the last six blocks



Orphan in Bitcoin (No reward)  
Ommers in Ethereum (Small reward)


Ommmer block defined as “The child of an ancestor that is not itself an ancestor”

- Smaller reward given for ommer blocks
- Transactions from ommer blocks are not included

## 2 Types of accounts in Ethereum

### Externally Owned Accounts (EOA)

Controlled by people + private keys



### Contract accounts (Smart Contracts)

Controlled by smart contract code + storage




Note that because a contract account does not have a private key, it cannot initiate a transaction. Only EOAs can initiate transactions, but contracts can react to transactions by calling other contracts, building complex execution paths.

# The Account State

The account state comprises the following four fields:

- balance: A scalar value equal to the number of Wei owned by this address.
- storageRoot: A 256-bit hash of the root node of a Merkle Patricia tree that encodes the storage contents of the account (a mapping between 256-bit integer values)
- codeHash: The hash of the EVM code of this account
- nonce: A scalar value equal to the number of transactions sent from this address or, in the case of accounts with associated code, the number of contract-creations made by this account.

# Ethereum Transactions





# State

14c5f8ba:  
- 1024 eth

bb75a980:  
- 5202 eth

```
if !contract.storage[tx.data[0]]:  
    contract.storage[tx.data[0]] = tx.data[1]
```

[0, 235235, 0, ALICE ...

892bf92f:  
- 0 eth

```
send(tx.value / 3, contract.storage[0])  
send(tx.value / 3, contract.storage[1])  
send(tx.value / 3, contract.storage[2])
```

[ALICE, BOB, CHARLIE]

4096ad65  
- 77 eth

## Transaction

From:  
14c5f8ba  
To:  
bb75a980  
Value:  
10  
Data:  
2,  
CHARLIE  
Sig:  
30452fdedb3d  
f7959f2ceb8a1

# State'

14c5f8ba:  
- 1014 eth

bb75a980:  
- 5212 eth

```
if !contract.storage[tx.data[0]]:  
    contract.storage[tx.data[0]] = tx.data[1]
```

[0, 235235, CHARLIE, ALICE ...

892bf92f:  
- 0 eth

```
send(tx.value / 3, contract.storage[0])  
send(tx.value / 3, contract.storage[1])  
send(tx.value / 3, contract.storage[2])
```

[ALICE, BOB, CHARLIE]

4096ad65  
- 77 eth

The Ethereum state transition function,  $\text{APPLY}(S, TX) \rightarrow S'$  can be defined as follows:

1. Check if the transaction is well-formed (ie. has the right number of values), the signature is valid, and the nonce matches the nonce in the sender's account. If not, return an error.
2. Calculate the transaction fee as  $\text{STARTGAS} * \text{GASPRICE}$ , and determine the sending address from the signature. Subtract the fee from the sender's account balance and increment the sender's nonce. If there is not enough balance to spend, return an error.
3. Initialize  $\text{GAS} = \text{STARTGAS}$ , and take off a certain quantity of gas per byte to pay for the bytes in the transaction.
4. Transfer the transaction value from the sender's account to the receiving account. If the receiving account does not yet exist, create it. If the receiving account is a contract, run the contract's code either to completion or until the execution runs out of gas.
5. If the value transfer failed because the sender did not have enough money, or the code execution ran out of gas, revert all state changes except the payment of the fees, and add the fees to the miner's account.
6. Otherwise, refund the fees for all remaining gas to the sender, and send the fees paid for gas consumed to the miner.

A submitted transaction includes the following information:

- `recipient` – the receiving address (if an externally-owned account, the transaction will transfer value. If a contract account, the transaction will execute the contract code)
- `signature` – the identifier of the sender. This is generated when the sender's private key signs the transaction and confirms the sender has authorised this transaction
- `value` – amount of ETH to transfer from sender to recipient (in WEI, a denomination of ETH)
- `data` – optional field to include arbitrary data
- `gasLimit` – the maximum amount of gas units that can be consumed by the transaction. Units of gas represent computational steps
- `maxPriorityFeePerGas` - the maximum amount of gas to be included as a tip to the miner
- `maxFeePerGas` - the maximum amount of gas willing to be paid for the transaction (inclusive of `baseFeePerGas` and `maxPriorityFeePerGas`)


## Some practical points about transactions

- Miners choose which transactions to include in a block
- Miners can add their own transactions to a block
- Miners choose the order of transactions in a block
- Your transaction is in competition with other transactions for inclusion in the block

# Processing a transaction

Initial account balance = 5

line 1  
line 2  
line 3  
line 4  
line 5



Process flow

Final account balance = 10

# Transactions are atomic

Initial account balance = 5

EVM reverts transaction

line 1  
line 2  
line 3  
line 4  
line 5

Final account balance = 5

Process flow



# Analysing Ethereum Data

# Block Explorer

## The Ethereum Blockchain Explorer

All Filters

Search by Address / Txn Hash / Block / Token / Ens



ETHER PRICE

\$3,345.60 @ 0.06645 BTC (+2.72%)



TRANSACTIONS

1,256.40 M (15.7 TPS)

MED GAS PRICE

39 Gwei (\$2.74)

ETHEREUM TRANSACTION HISTORY IN 14 DAYS



MARKET CAP

\$392,180,105,051.00



DIFFICULTY

8,075.63 TH

HASH RATE

616,077.75 GH/s



### Latest Blocks

Bk	13081315 27 secs ago	Miner firepool 257 txns in 20 secs	3,71611 Eth
Bk	13081313 52 secs ago	Miner BeePool 65 txns in 26 secs	2,03172 Eth
Bk	13081312 1 min 18 secs ago	Miner 0x0708f87a089a91c65d...	2,17696 Eth
Bk	13081311 1 min 50 secs ago	Miner Ethermine 383 txns in 8 secs	2,2938 Eth
Bk	13081310 1 min 58 secs ago	Miner Nanopool 251 txns in 2 secs	2,2438 Eth
Bk	13081309 2 min 0 secs ago	Miner BeePool 82 txns in 10 secs	2,03181 Eth

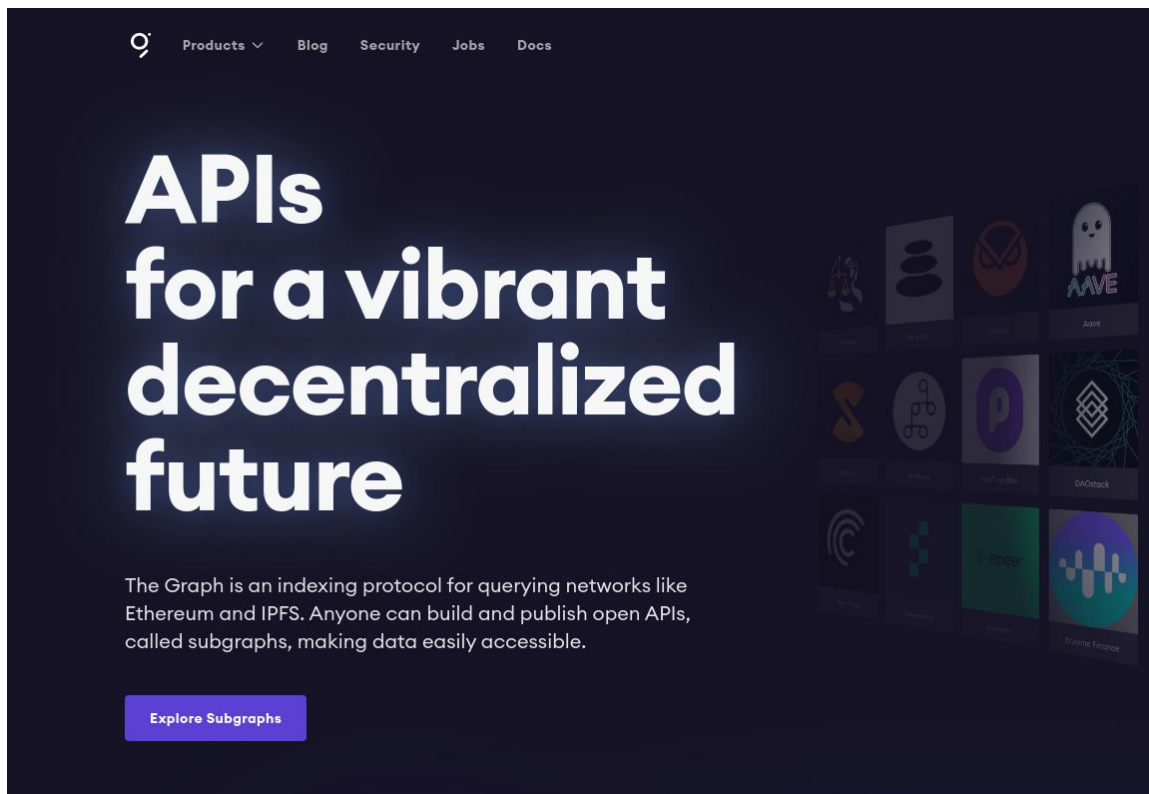
[View all blocks](#)

### Latest Transactions

Tx	0x9f6520ded15e...	From 0x8a689700adbfc2dab7... To 0x198fa2c6f9591ac9fd9...	0.00556 Eth
Tx	0xcb79e81bc04bf...	From 0xcdfa1be78735b56ad9... To 0x2f151af3999d339f3a9...	0.0006 Eth
Tx	0x3cf1e4b736e...	From 0x5d9d4a102bcf2d8ed8... To 0x7a250d5630b4cf5397...	0 Eth
Tx	0x297d2f1d02cf0...	From 0x92573646c0c351f666... To 0x1a2a1c938ce3ec39b6...	0 Eth
Tx	0xa275cfdea754...	From 0xffbcd95ef09e297dc33d... To 0x31c8eacbfdd875c74b...	0 Eth
Tx	0xa4534caf765a...	From 0x2d5427f30e0085d25c... To 0x7a7a75d4f6a30b4e4f5307...	0.2 Eth

[View all transactions](#)





The screenshot shows the homepage of The Graph. At the top, there is a navigation bar with a search icon, a 'Products' dropdown menu, and links for 'Blog', 'Security', 'Jobs', and 'Docs'. The main heading is 'APIs for a vibrant decentralized future' in large white text. Below this, a paragraph explains that The Graph is an indexing protocol for querying networks like Ethereum and IPFS, allowing anyone to build and publish open APIs called subgraphs. A purple button labeled 'Explore Subgraphs' is positioned below the text. On the right side, there is a grid of 16 subgraph icons, each with a unique logo and name, such as 'Aave', 'Uniswap', and 'Chainlink'.

Products ▾ Blog Security Jobs Docs

# APIs for a vibrant decentralized future

The Graph is an indexing protocol for querying networks like Ethereum and IPFS. Anyone can build and publish open APIs, called subgraphs, making data easily accessible.

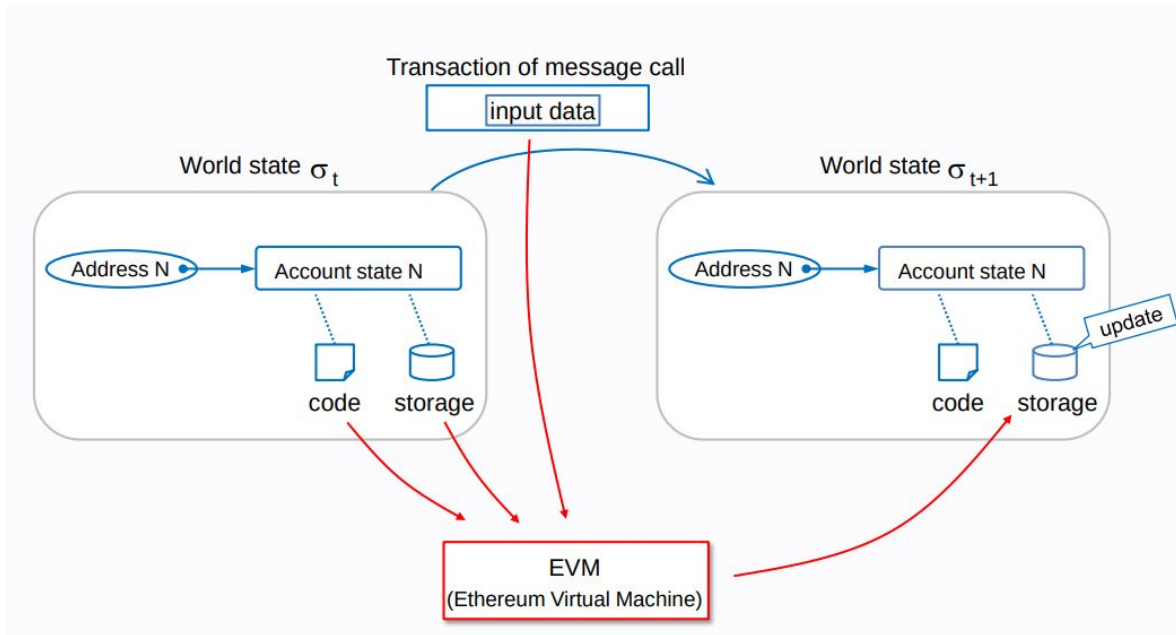
[Explore Subgraphs](#)

The Graph

# Ethereum Virtual Machine

Ethereum is a distributed state machine.

Ethereum's state is a large data structure which holds not only all accounts and balances, but a *machine state*, which can change from block to block according to a pre-defined set of rules, and which can execute arbitrary machine code.



From : [Ethereum EVM](#)

The EVM is a stack machine , the stack has a maximum size of 1024.

Stack items have a size of 256 bits; in fact, the EVM is a 256-bit word machine (this facilitates Keccak256 hash scheme and elliptic-curve computations).

During execution 2 areas are available for variables

- memory - a transient memory which does not persist between transactions
- storage - part of a Merkle Patricia storage trie associated with the contract's account, part of the global state

# Execution / OP codes

A compiled smart contract bytecode executes as a number of EVM opcodes

There is a maximum of 256 possible op codes : [Reference sheet](#) for op codes

Op codes can be categorised as

- **Stack-manipulating opcodes** (*POP, PUSH, DUP, SWAP*)
- **Arithmetic/comparison/bitwise opcodes** (*ADD, SUB, GT, LT, AND, OR*)
- **Environmental opcodes** (*CALLER, CALLVALUE, NUMBER*)
- **Memory-manipulating opcodes** (*MLOAD, MSTORE, MSTORE8, MSIZE*)
- **Storage-manipulating opcodes** (*SLOAD, SSTORE*)
- **Program counter related opcodes** (*JUMP, JUMPI, PC, JUMPDEST*)
- **Halting opcodes** (*STOP, RETURN, REVERT, INVALID, SELFDESTRUCT*)

# EVM Languages

See : <https://ethereum.org/en/developers/docs/smart-contracts/languages/>

- Solidity
  - The most popular programming language for Ethereum contracts
- LLL
  - *Low-level Lisp-like Language*
- Vyper
  - A language with overflow-checking, numeric units but without unlimited loops
- Yul / Yul+
  - An intermediate language that can be compiled to bytecode for different backends. Support for EVM 1.0, EVM 1.5 and Ewasm is planned, and it is designed to be a usable common denominator of all three platforms.
- FE
  - Statically typed language Inspired by Rust and Python
- Pyramid Scheme (experimental)
  - A Scheme compiler into EVM that follows the SICP compilation approach
- Flint
  - A language with several security features: e.g. asset types with a restricted set of atomic operations
- LLLL
  - An LLL-like compiler being implemented in Isabelle/HOL
- HAssembly-evm
  - An EVM assembly implemented as a Haskell DSL
- Bamboo (experimental)
  - A language without loops

## Vyper

- Pythonic programming language
- Strong typing
- Small and understandable compiler code
- Deliberately has less features than Solidity with the aim of making contracts more secure and easier to audit. Vyper

does not support:

- Modifiers
- Inheritance
- Inline assembly
- Function overloading
- Operator overloading
- Recursive calling
- Infinite-length loops
- Binary fixed points

## Example Vyper Code

```
beneficiary: public(address)
auctionStart: public(uint256)
auctionEnd: public(uint256)

pendingReturns: public(HashMap[address, uint256])

# Create a simple auction with `_bidding_time`
@external
def __init__(_beneficiary: address, _bidding_time: uint256):
    self.beneficiary = _beneficiary
    self.auctionStart = block.timestamp
    self.auctionEnd = self.auctionStart + _bidding_time
```

```
@external
@payable
def bid():

    assert block.timestamp <
self.auctionEnd
    # Check if bid is high enough
    assert msg.value > self.highestBid
self.pendingReturns[self.highestBidder]
+= self.highestBid
    # Track new high bid
self.highestBidder = msg.sender
self.highestBid = msg.value
```



## Yul

- Intermediate language for Ethereum.
- Supports the [EVM](#) and [Ewasm](#), an Ethereum flavored WebAssembly, and is designed to be a usable common denominator of both platforms.
- Good target for high-level optimisation stages that can benefit both EVM and Ewasm platforms equally.

## Yul+

- A low-level, highly efficient extension to Yul.
- Initially designed for an [optimistic rollup](#) contract.
- Yul+ can be looked at as an experimental upgrade proposal to Yul, adding new features to it.

## Yul Example

```
{
  function power(base, exponent) -> result
  {
    switch exponent
    case 0 { result := 1 }
    case 1 { result := base }
    default
    {
      result := power(mul(base, base), div(exponent, 2))
      if mod(exponent, 2) { result := mul(base, result) }
    }
  }
  let res := power(calldataload(0), calldataload(32))
  mstore(0, res)
  return(0, 32)
}
```

# FE

See : <https://ethereum.org/en/developers/docs/smart-contracts/languages/#fe>

- Statically typed language for the Ethereum Virtual Machine (EVM).
- Inspired by Python and Rust.
- Aims to be easy to learn – even for developers who are new to the Ethereum ecosystem.
- Fe development is still in its early stages, the language had its alpha release in January 2021.

## FE Example

```
type BookMsg = bytes[100]
```

```
contract GuestBook:
```

```
    pub guest_book: map<address, BookMsg>
```

```
    event Signed:
```

```
        book_msg: BookMsg
```

```
    pub def sign(book_msg: BookMsg):
```

```
        self.guest_book[msg.sender] = book_msg
```

```
        emit Signed(book_msg=book_msg)
```

```
    pub def get_msg(addr: address) -> BookMsg:
```

```
        return self.guest_book[addr].to_mem()
```

## Comparison / Cheat Sheet for Solidity and Vyper

<https://reference.auditless.com/cheatsheet/>

# Ethereum Clients

Name	Language	Sync Strategy	State Pruning
Geth	Go	Fast, Full	Archive, Pruned
Open Ethereum	Rust	Warp, Full	Archive, Pruned
Nethermind	C#	Fast, Full	Archive, Pruned
Besu	Java	Fast, Full	Archive, Pruned
Erigon	Go	Fast, Full	Archive, Pruned

## Sync Modes

- Full – downloads all blocks (including headers, transactions and receipts) and generates the state of the blockchain incrementally by executing every block.
- Fast (Default) – downloads all blocks (including headers, transactions and receipts), verifies all headers, and downloads the state and verifies it against the headers.
- Light – downloads all block headers, block data, and verifies some randomly.
- Warp sync – Every 5,000 blocks, nodes will take a consensus-critical snapshot of that block's state. Any node can fetch these snapshots over the network, enabling a fast sync.
- Beam sync – A sync mode that allows you to get going faster. It doesn't require long waits to sync, instead it back-fills data over time.
- Header sync – you can use a trusted checkpoint to start syncing from a more recent header and then leave it up to a background process to fill the gaps eventually

# Mining on Ethereum

Miners run Ethash (PoW algorithm)

- Uses a 4 GB dataset, which is updated every 30,000 blocks

Ethash is intended to satisfy the following goals:

1. IO saturation: The algorithm should consume nearly the entire available memory access bandwidth (this is a strategy toward achieving ASIC resistance)
2. GPU friendliness: We try to make it as easy as possible to mine with GPUs. Targeting CPUs is almost certainly impossible, as potential specialization gains are too great, and there do exist criticisms of CPU-friendly algorithms that they are vulnerable to botnets, so we target GPUs as a compromise.
3. Light client verifiability: a light client should be able to verify a round of mining in under 0.01 seconds on a desktop in C, and under 0.1 seconds in Python or Javascript, with at most 1 MB of memory (but exponentially increasing)
4. Light client slowdown: the process of running the algorithm with a light client should be much slower than the process with a full client, to the point that the light client algorithm is not an economically viable route toward making a mining implementation, including via specialized hardware.
5. Light client fast startup: a light client should be able to become fully operational and able to verify blocks within 40 seconds in Javascript.

From [Ethereum Wiki](#)



# Ethereum Test Networks

The Ropsten test network is a Proof-of-Work testnet for Ethereum. To acquire ETH on Ropsten, one can mine on the network.

The Kovan test network is a Proof-of-Authority testnet for Ethereum, originally started by the Parity team. To acquire ETH on Kovan, one can request it from a faucet.

The Rinkeby test network is a Proof-of-Authority testnet for Ethereum, originally started by the Geth team. To acquire ETH on Rinkeby, one can request it from a faucet.

The Görli test network is a Proof-of-Authority testnet for Ethereum, originally proposed by Chainsafe and Afri Schoedon. To acquire ETH on Görli, one can use the one-way throttled bridge from any of the other three test networks.